


How the agent works — the self-improving goal loop

One fixed model (Opus 4.8), three separated roles. The model never gets smarter — **the files around it do.**



From idea to shipped — the command sequence

Everything runs from a session opened in the agent home. Your taste and sign-offs are the only manual steps — by design.

- 1** **/prd <product idea>**
or /deep-prd to force exhaustive, verified research
Researches best practices + analog software, mines the **requirements you never stated**, writes the PRD + a seeded rubric. **You:** answer up to 4 scope questions.
- 2** **/design-direction <product>, target: <path>**
only for products with a UI — run BEFORE the first UI slice
Interactive taste session → frozen tokens, reference images, banned visual outcomes. **You:** bring screenshots of UIs you admire; react live; pick a direction.
- 3** **/goal-opus implement phase 1 of <PRD>, target: <path>**
one roadmap phase = one goal run, using criteria.seed.json
The loop from slide 1 runs to done or honest abort. **You: sign off the rubric — this is your steering wheel.** Shape criteria, banned outcomes, max iterations.
- 4** **repeat per phase · demo each slice**
phase gates must pass before the next phase opens
At slice demos, give the feedback screenshots can't capture: **motion feel, interactive performance**. That's the one thing the loop honestly can't judge.
- 5** **read the final report**
criteria table · iterations used · what was learned & written back where
Trust an **aborted** verdict — it means the system refused to fake success. Everything learned is already committed; the next run starts smarter.

Rules of thumb — using it efficiently

The habits that keep the system compounding instead of drifting.

One home, many targets

- ▶ **Never clone this repo per project** — clones fork the self-improving skill and splinter its learning.
- ▶ Every build goal declares **target: <path>**; product code + its own STATE.md live there.
- ▶ The home keeps only the system, memory, and run evidence.

Where memory lives

- ▶ **STATE.md** — verified facts, decisions D1–D8, open failures, resume pointer.
- ▶ **SKILL.md sections** — known failure modes, anti-patterns, eval suite, run log (append-only, dated).
- ▶ **git history** — every self-edit is a diff you can audit and revert.

Your control points

- ▶ **Rubric sign-off** (Phase 1) — the main steering wheel.
- ▶ **Design direction** — taste enters here, once, then freezes.
- ▶ **Slice demos** — motion & feel feedback.
- ▶ Everything else runs hands-off.

Research depth

- ▶ **/prd** judges for itself when to escalate to deep research — and says which depth it chose.
- ▶ **/deep-prd** forces it, and fails loudly if the plugin is missing — never a silent downgrade.

Settled — don't re-litigate

- ▶ **No RAG / vector index; no maker-side code-graph** — grep over live files won on evidence (0.92 vs 0.83); details in docs/context/11.
- ▶ **No enforcement hooks** — they can veto work the rubric requires.
- ▶ Minimalism ladder + completeness guard are live in every run.

Gotchas

- ▶ New skills/agents register at **session start** — restart after installing, retry once if an agent isn't found.
- ▶ Goals into a fresh session, not the one that scaffolded the files.
- ▶ First UI slice E2E-tests vision-verify — check the verdict names the image files it viewed.